




Multi-way Search Trees

- ⊕ Each internal node of a multi-way search tree T :
 - ⊕ has at least two children
 - ⊕ contains $d - 1$ items, where d is the number of children → d -nodes
 - ⊕ “contains” 2 pseudo-items: $k_0 = -\infty$, $k_d = \infty$
- ⊕ Children of each internal node are “between” items
- ⊕ all keys in the subtree rooted at the child fall between keys of those items



Multi-way Searching

- Similar to binary searching
 - If search key $s < k_1$ search the leftmost child
 - If $s > k_{d-1}$, search the rightmost child
- That's it in a binary tree; what about if $d > 2$?
 - Find two keys k_{i-1} and k_i between which s falls, and search the child v_i .
- What would an in-order traversal look like?






2-4 Trees

- a. Nodes may contain 1, 2 or 3 items.
- b. A node with k items has $k + 1$ children
- c. All leaves are on same level.



Example






Insertion

- ✚ Insertion:
- ✚ Find the appropriate leaf. If there is only one or two items, just add to leaf.
- ✚ If no room, move middle item to parent and split remaining two items among two children.



Insertion

insert 80




Overflow!



Insertion

Split & move middle element to parent





Removal

- ⊕ First : find the key with a simple multi-way search
- ⊕ If the item to delete is in an internal node, reduce to the case where item is at the bottom of the tree by:
 - ⊠ Find item which precedes it in in-order traversal
 - which one?
 - ⊠ Swap them
 - ⊠ Remove the item
 - ⊠ **Alternative?**



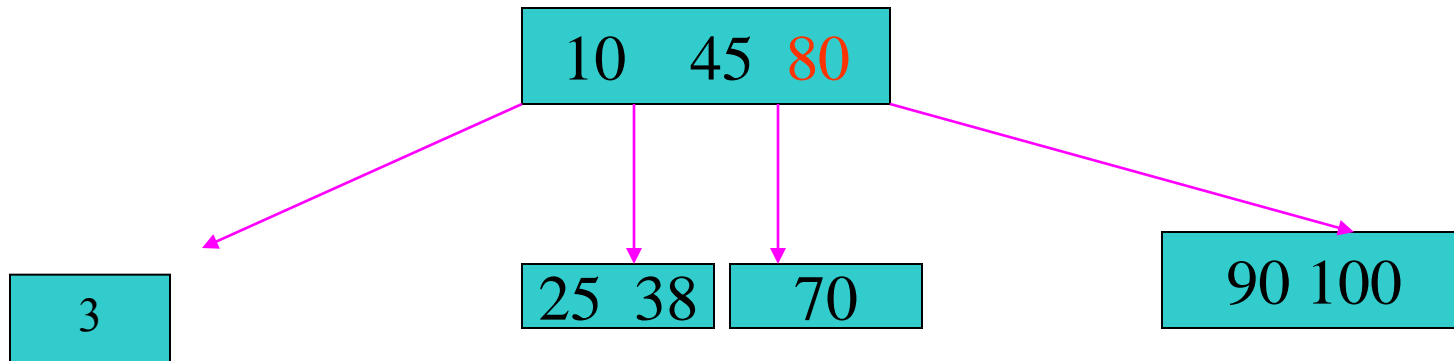
Removal

- ⊕ Not enough items in the node
Underflow!
- ⊕ Pull an item from the parent, replace it with an item from a sibling - transfer
- ⊕ Still not good enough! What happens if siblings are 2-nodes?
- ⊕ Could we just pull one item from the parent?



Removal


- Remove 3
 - move 10 into the subtree
 - move 25 into the parent





Removal


- ✱ If siblings are 2-nodes (i.e. contain only one key)
 - ✱ cannot 'steal' from them
- ✱ Do node merging
- ✱ Remove 3
 - ✱ move 10 into the subtree
 - ✱ merge 10 with 25





2-4 Trees

- More on removal:
 - What if parent is a 2-node?
 - Propagate underflow up the tree
 - Delete 9





2-4 Trees

- ⊕ 2-4 trees are easy to maintain
- ⊕ Insertion and deletion take $O(\log n)$
- ⊕ Balanced trees



B-Trees

- ⊙ Up to now, all data that has been stored in the tree has been in memory.
- ⊙ If data gets too big for main memory, what do we do?
- ⊙ If we keep a pointer to the tree in main memory, we could bring in just the nodes that we need.
- ⊙ For instance, to do an insert with a BST, if we need the left child, we do a disk access and retrieve the left child.
- ⊙ If the left child is NIL, then we can do the insert, and store the child node on the disk.
- ⊙ Not too good for a BST



B-Trees

- ⊙ The problem with BST: storing the data requires disk accesses, which is expensive, compared to execution of machine instructions.
- ⊙ If we can reduce the number of disk accesses, then the procedures run faster.
- ⊙ The only way to reduce the number of disk accesses is to increase the number of keys in a node.
- ⊙ The BST allows only one key per leaf.
- ⊙ **Very good and often used for Search Engines!**
 - (when collection size gets very big → the index does not fit in memory)




B-Trees

- If we increase the number of keys in the nodes, how will we do any tree operations effectively?

10 20 30 40 50 60 70


- Above is a node with 7 keys. How do we add children?





B-Trees


- Clearly, the tree below is useless.



- How many pointers do we need?
- Using the idea of BST, we need to be able to put nodes into the tree that have smaller, same and larger values than the node we are currently examining.



B-Trees: A General Case of Multi-Way Search Trees



- We can easily find any value.
- We need to create operations, which require rules on what makes a tree a B-Tree.
- Clearly, having one key per node would be very bad.
- We need a mechanism to increase the height of the tree (since the number of keys in any node can get very high) so we can shift keys out of a node, making the nodes smaller.



B-Trees: Fields in a Node

A B-Tree is a rooted tree (whose root is $\text{root}[T]$)
having the following properties:

1. Every internal node x has the following fields:

$n[x]$	$\text{key}_1[x]$]	$\text{key}_2[x]$]	$\text{key}_3[x]$]	$\text{key}_4[x]$]	$\text{key}_5[x]$]	$\text{key}_6[x]$]	$\text{key}_7[x]$]	$\text{key}_8[x]$]	$\text{leaf}[x]$]
$c_1[x]$	$c_2[x]$	$c_3[x]$	$c_4[x]$	$c_5[x]$	$c_6[x]$	$c_7[x]$	$c_8[x]$	$c_9[x]$	

$n[x]$ is the number of keys in the node. $n[x] = 8$ above.

$\text{leaf}[x] = \text{false}$ for internal nodes, since x is not a leaf.

The $\text{key}_i[x]$ are the values of the keys, where $\text{key}_i[x] \leq \text{key}_{i+1}[x]$.

$c_i[x]$ are pointers to child nodes. All the keys in $c_i[x]$ have values that are between $\text{key}_{i-1}[x]$ and $\text{key}_i[x]$.



B-Trees

- ⊕ Leaf nodes have no child pointers
- ⊕ $\text{leaf}[x] = \text{true}$ for leaf nodes.
- ⊕ All leaf nodes are at the same level

$n[x]$	$\text{key}_1[x]$]	$\text{key}_2[x]$]	$\text{key}_3[x]$]	$\text{key}_4[x]$]	$\text{key}_5[x]$]	$\text{key}_6[x]$]	$\text{key}_7[x]$]	$\text{key}_8[x]$]	$\text{leaf}[x]$]
--------	------------------------	------------------------	------------------------	------------------------	------------------------	------------------------	------------------------	------------------------	-----------------------



B-Trees

- ✚ There are lower and upper bounds on the number of keys a node can contain. This depends on the “**minimum degree**” $t \geq 2$, which we must specify for any given B-Tree.
- a. Every node other than the root must have at least $t-1$ keys. Every internal node other than the root thus has t children. If the tree is nonempty, the root must have at least one key.
- b. Every node can contain at most $2t-1$ keys. Therefore, an internal node can have at most $2t$ children. A node is **full** if it contains exactly $2t-1$ keys.



Height of B-Tree

- ✚ If $n \geq 1$, then for any n -key B-tree of height h and minimum degree $t \geq 2$,

$$\text{height} = h \leq \log_t[(n+1)/2]$$

- ✚ The important thing to notice is that the height of the tree is log base t . So, as t increases, the height, for any number of nodes n , will decrease.
- ✚ Using the formula $\log_a x = (\log_b x)/(\log_b a)$, we can see that
 - ✚ $\log_2 10^6 = (\log_{10} 10^6)/(\log_{10} 2) \approx 6/0.30102999566398 \approx 19$
 - ✚ $\log_{10} 10^6 = 6$

So, 13 less disk accesses to get to the leafs!



Basic Operations

- ❖ The root of the B-tree is always in main memory, so that a Disk-Read on the root is never required; a Disk-Write of the root is required, however, whenever the root node is changed.
- ❖ Any nodes that are passed as parameters must already have had a Disk-Read operation performed on them.



Searching a B-Tree

B-TREE-SEARCH(x, k)

```
1   $i \leftarrow 1$ 
2  while  $i \leq n[x]$  and  $k > key_i[x]$ 
3      do  $i \leftarrow i + 1$ 
4  if  $i \leq n[x]$  and  $k = key_i[x]$ 
5      then return  $(x, i)$ 
6  if  $leaf[x]$ 
7      then return NIL
8  else DISK-READ( $c_i[x]$ )
9      return B-TREE-SEARCH( $c_i[x], k$ )
```

Start at the leftmost key in the node, and go to the right until you go too far.

If it is a leaf node, then you are done, as there is no leaf to inspect

Otherwise, retrieve the child node from the disk, and put it into memory



Inserting into B-trees

- ✚ Really very easy. Very similar with (2,4) trees.
- ✚ Just keep in mind that you are starting at the root, and then finding the subtree where the key should be inserted, and following the pointer.
- ✚ A deletion may eventually occur, and sometimes deletions force keys into their parents. So, if we encounter a full node on our way to the node where the insertion will take place, we must split that node into two.



Inserting into B-trees (cont'd)

B-TREE-INSERT(T, k)

```
1   $r \leftarrow \text{root}[T]$ 
2  if  $n[r] = 2t - 1$ 
3      then  $s \leftarrow \text{ALLOCATE-NODE}()$ 
4           $\text{root}[T] \leftarrow s$ 
5           $\text{leaf}[s] \leftarrow \text{FALSE}$ 
6           $n[s] \leftarrow 0$ 
7           $c_1[s] \leftarrow r$ 
8          B-TREE-SPLIT-CHILD( $s, 1, r$ )
9          B-TREE-INSERT-NONFULL( $s, k$ )
10 else B-TREE-INSERT-NONFULL( $r, k$ )
```

If the node has $2t-1$ keys, it can't accept any more keys, so you need to split it into 2 nodes before doing the insert.

Otherwise, call Nonfull()



Deleting Keys from Nodes

1. If the key k is in node x and x is a leaf, delete the key k from x .
2. If the key k is in node x and x is an internal node, do the following.
 - a. If the child y that precedes k in node x has at least t keys, then find the predecessor k' of k in the subtree rooted at y . Recursively delete k' , and replace k by k' in x . (Finding k' and deleting it can be performed in a single downward pass.)
 - b. Symmetrically, if the child z that follows k in node x has at least t keys, then find the successor k' of k in the subtree rooted at z . Recursively delete k' , and replace k by k' in x . (Finding k' and deleting it can be performed in a single downward pass.)
 - c. Otherwise, if both y and z have only $t - 1$ keys, merge k and all of z into y , so that x loses both k and the pointer to z , and y now contains $2t - 1$ keys. Then, free z and recursively delete k from y .



Deleting Keys from Nodes

3. If the key k is not present in internal node x , determine the root $c_i[x]$ of the appropriate subtree that must contain k , if k is in the tree at all. If $c_i[x]$ has only $t - 1$ keys, execute step 3a or 3b as necessary to guarantee that we descend to a node containing at least t keys. Then, finish by recursing on the appropriate child of x .
 - a. If $c_i[x]$ has only $t - 1$ keys but has an immediate sibling with at least t keys, give $c_i[x]$ an extra key by moving a key from x down into $c_i[x]$, moving a key from $c_i[x]$'s immediate left or right sibling up into x , and moving the appropriate child pointer from the sibling into $c_i[x]$.
 - b. If $c_i[x]$ and both of $c_i[x]$'s immediate siblings have $t - 1$ keys, merge $c_i[x]$ with one sibling, which involves moving a key from x down into the new merged node to become the median key for that node.